TxShield: Transaction-level RAG-enabled DeFi Attack Detection

Wei Ma†, Shiya Li†, Jiayi Xu, Ye Liu*, Junjie Shi, Cong Wu, Jing Chen, Lingxiao Jiang, and Zijian Zhang

***Abstract*—Decentralized Finance (DeFi) protocols have become prime targets for sophisticated attacks, with losses exceeding $3.4 billion in 2025. However, existing detection approaches face fundamental limitations. Traditional rule-based detection tools are limited in their ability to capture diverse attack behaviors at runtime, while advanced large language models (LLMs) exhibit limitations due to insufficient task-specific knowledge and hallucinations. To this end, we observe that DeFi attacks, despite diverse implementations, exhibit recurring behavioral patterns at the semantic level, suggesting that their apparent diversity masks a small number of underlying semantic structures. In this paper, we present TxShield, a Retrieval-Augmented Generation framework for transaction-level DeFi attack detection. TxShield constructs a knowledge base of historical attacks with LLM-generated behavioral summaries, retrieves historical cases and applies similarity filtering to select relevant historical context, and selects prompts according to the available context to guide LLM reasoning by analogy with retrieved patterns. XGBoost estimates transaction risk to control the use of retrieved cases and provides an auxiliary probability hint for LLM reasoning. The LLM produces the final classification.**

On an evaluation set constructed from 364 DeFiHackLabs attack transactions and 364 benign transactions, TxShield achieves 92.93% F1 and 97.87% precision. Compared with direct XGBoost classification, it improves F1 by 6.45 percentage points and reduces the number of false-positive transactions from 60 to 7.

***Index Terms*—Smart Contract Security, DeFi Attack Detection, Retrieval-Augmented Generation, Transaction Analysis, Blockchain Security**

I. Introduction

Decentralized Finance (DeFi) has emerged as a major application of blockchain technology, enabling permissionless financial services through smart contracts. By late 2025, the total value locked (TVL) in DeFi protocols had exceeded $230 billion, reflecting the substantial financial value managed by DeFi systems [1]. This rapid growth has attracted a large user base, with tens of millions of users interacting with DeFi protocols worldwide. The increasing scale of DeFi has been accompanied by significant security challenges. Industry analyses report that losses from cryptocurrency hacks and exploits exceeded $3.4 billion in 2025, with DeFi protocols representing a major attack surface [2], [3]. These attacks span both traditional smart-contract vulnerabilities, such as reentrancy as seen in the DAO incident, and advanced economic exploits that leverage flash loans to manipulate prices and protocol states within a single transaction [4]. The immutable nature of smart contracts exacerbates these risks. Once deployed, vulnerabilities cannot be patched retroactively, making effective testing and continuous post-deployment monitoring critical for mitigating security threats in DeFi systems.

Existing approaches to smart contract security can be broadly categorized into four paradigms, each with inherent limitations. *Static analysis tools* such as Slither [5] and Mythril [6] use pattern matching and symbolic execution to identify known vulnerability patterns. While effective at detecting classic bugs such as reentrancy, these tools suffer from path explosion in complex contracts and often struggle to capture runtime behaviors or cross-protocol interactions arising from DeFi composability [7]. *Formal verification frameworks* like Certora Prover [8] and KEVM [9] provide mathematical guarantees but require substantial manual effort to specify invariants and often struggle to model the dynamic, multi-protocol nature of DeFi. *Rule-based dynamic detection tools* such as DeFiRanger [10], DeFiTainter [11], and DeFort [12] analyze transaction traces to detect price-manipulation attacks. However, these approaches rely on predefined mathematical models of specific pricing mechanisms (e.g., constant-product AMMs), making them less effective for custom protocols with non-standard price models; DeFiScope reports that such custom price models account for 44.2% of 95 real-world price manipulation attacks collected over the past three years [13]. *LLM-based approaches* represent a recent paradigm shift. TrxGNNBERT [14] leverages a graph-transformer language model to learn trace-level representations for DeFi fraud detection, while GPTScan [15] combines GPT with static analysis for logic-vulnerability detection and iAudit [16] employs fine-tuned models with multi-agent debate for source-code auditing. However, these methods either focus primarily on pre-deployment source-code analysis and cannot detect runtime attacks, or incur high end-to-end overhead for trace construction and model inference, limiting practicality for real-time monitoring. LLM hallucinations can lead to plausible but incorrect vulnerability reports. Moreover, parametric knowledge may lack the transaction-specific context needed to distinguish an attack from benign behavior.

Our key insight is that DeFi attacks, despite their diverse manifestations, often exhibit recurring behavioral patterns at the semantic level. For example, a flash-loan-based price

†Wei Ma and Shiya Li contributed equally to this work and share first authorship. *Corresponding author: Ye Liu.

Wei Ma is an Assistant Professor at Blekinge Institute of Technology (BTH), Sweden. This work was mainly carried out while he was at Singapore Management University (SMU) (e-mail: weima@smu.edu.sg).

Lingxiao Jiang is with Singapore Management University.

Shiya Li, Jiayi Xu, Cong Wu, and Jing Chen are with Wuhan University (e-mail: 2022302181203@whu.edu.cn; 2022302181176@whu.edu.cn; cnacwu@whu.edu.cn; chenjing@whu.edu.cn).

Ye Liu and Zijian Zhang are with Beijing Institute of Technology (e-mail: ye.liu@bit.edu.cn; zhangzijian@bit.edu.cn).

Junjie Shi is with Nanyang Technological University (e-mail: chunkit001@e.ntu.edu.sg).

manipulation attack can follow a "borrow, manipulate, and profit" sequence: obtaining large capital via flash loans, executing trades that artificially inflate or deflate token prices, and extracting profit before repaying the loan, all within a single atomic transaction. This observation suggests that historical attack cases, annotated with behavioral semantics, can serve as a valuable knowledge source to guide the detection of new attacks. To empirically ground this insight, we analyze approximately 200 real-world DeFi attack incidents from DeFiHackLabs [17]. Despite their apparent diversity, these attacks consistently fall into eight major semantic categories, with Business Logic Flaws, Price Manipulation, and Reentrancy-based call attacks dominating the incident landscape. These results provide quantitative evidence that attack patterns are reusable across different protocols.

However, realizing this insight poses three technical challenges: **(C1) Semantic Gap:** Raw execution traces consist of low-level EVM operations and function calls that are difficult to interpret. *How can we bridge the gap between low-level traces and high-level attack semantics?* **(C2) Retrieval Precision:** Naïve similarity search may retrieve irrelevant or misleading historical cases, introducing noise that can degrade detection accuracy. *How can we ensure that retrieved cases are truly relevant to the query transaction?* **(C3) Reasoning Reliability:** Even with relevant context, LLMs may hallucinate or fail to correctly apply retrieved knowledge. *How can we guide the LLM to make reliable attack-or-benign decisions?*

To address these challenges, we propose TxShield, a Retrieval-Augmented Generation framework for the DeFi transaction attack detection. At a high level, TxShield combines a *Historical Knowledge Base* with a *RAG-based Decision Engine* that retrieves cases, selects context, constructs prompts, and produces LLM judgments. For **C1**, we construct a knowledge base by collecting execution traces from historical attack incidents, extracting behavioral summaries using LLM-based semantic summarization, and storing them as vector embeddings alongside structured metadata. This transforms opaque execution traces into interpretable attack patterns. For **C2**, we design a two-stage retrieval mechanism: the first stage performs vector similarity search to retrieve top-k candidate cases; the second stage applies a similarity threshold ($\tau = 0.85$) to filter out low-confidence matches, retaining cases that meet the similarity criterion. For **C3**, XGBoost estimates risk from numerical transaction features before LLM inference. When retrieved cases are available, this risk estimate determines whether they are retained as context. The prompt also includes the attack probability as an auxiliary signal, while the LLM makes the final decision from transaction evidence. Transactions without retained context are analyzed without historical examples.

We construct DeFiAttackBench, an evaluation set combining 364 attack transactions from DeFiHackLabs [17] with 364 benign transactions sampled from the affected protocols. The system achieves 92.93% F1, compared with 86.48% for XGBoost-only, and reduces the number of false-positive transactions from 60 to 7. The variant without XGBoost achieves 77.94% F1. The backbone ablation examines detection performance across nine LLMs. On a 91-attack subset of the D1 price-manipulation dataset introduced by DeFiScope [13], which we denote PriceAttackBench, the system detects 86 attacks (94.51%). These experiments examine both general transaction detection and specialized price-manipulation detection.

This paper makes the following contributions:

- We present TxShield, the first RAG-based framework for the DeFi transaction attack detection that leverages historical attack knowledge to enhance LLM reasoning, addressing the limitations of both rule-based and pure LLM approaches.
- We introduce risk-guided context selection and prompting to address the ambiguity between resemblance to historical attacks and malicious intent. An XGBoost risk estimate controls the use of retrieved cases and informs LLM reasoning, while transaction evidence determines the final classification.
- We evaluate detection on 728 transactions, including a direct XGBoost baseline, component ablations, and backbone ablations across nine LLMs. The combined system improves F1 and reduces false positives relative to direct numerical classification.
- We characterize the strengths and limitations of trace-based semantic reasoning across attack categories, identifying the behavioral evidence that supports detection and the contextual requirements that limit it.

Roadmap. Section II introduces the background. Section III details the design of TxShield, including knowledge base construction, retrieval with similarity filtering, and XGBoost-guided retrieval and prompting. Section VI presents our experimental evaluation and ablation studies, Section VII discusses related work and Section VIII concludes with limitations and future directions.

II. Background

*a) **Smart Contracts and DeFi Ecosystem**:* Smart contracts are self-executing programs deployed on blockchain platforms such as Ethereum. Written primarily in Solidity, these contracts are compiled into bytecode and executed by the Ethereum Virtual Machine (EVM), a stack-based virtual machine that processes transactions deterministically across all network nodes. Once deployed, smart contract code is immutable (it cannot be modified or patched), which makes security vulnerabilities particularly consequential. Decentralized Finance (DeFi) leverages smart contracts to replicate traditional financial services without centralized intermediaries [7]. The DeFi ecosystem comprises several key protocol types: (1) *Decentralized Exchanges (DEXs)* such as Uniswap [18] enable permissionless token trading through Automated Market Makers (AMMs), where liquidity pools and mathematical formulas (e.g., the constant-product formula $x \cdot y = k$) determine token prices; (2) *Lending Protocols* like Aave [19] and Compound allow users to lend and borrow assets with interest rates determined algorithmically; and (3) *Flash Loan Providers* offer uncollateralized loans that must be borrowed and repaid within a single atomic transaction [4]. A defining characteristic of DeFi is *composability* (the ability for protocols to interact seamlessly, with outputs from one protocol serving as inputs

to another). While composability enables capital efficiency and innovation, it also creates complex attack surfaces where vulnerabilities can span multiple protocols.

*b) **DeFi Attack Taxonomy**:* Based on existing literature [7], [20], [21], we categorize DeFi attacks into the following major types:

Reentrancy Attacks. Reentrancy occurs when a contract makes an external call to another contract before updating its own state, allowing the callee to "re-enter" the original function and exploit the inconsistent state [22]. Modern variants include cross-function reentrancy (exploiting state shared across multiple functions) and cross-contract reentrancy (exploiting interactions between multiple contracts) [23], [24].

Price Manipulation Attacks. These attacks exploit the mechanism by which DeFi protocols determine asset prices. In AMM-based DEXs, an attacker can execute large trades to temporarily shift the price curve, then exploit this artificial price in another protocol (e.g., a lending platform using the DEX as a price oracle) [10], [25]. Flash loans amplify this threat by providing attackers with massive capital at zero upfront cost [4]. Time-Weighted Average Price (TWAP) oracles, designed to mitigate single-block manipulation, remain vulnerable to multi-block attacks [26].

Access Control Vulnerabilities. These vulnerabilities arise when privileged functions lack proper authorization checks, allowing unauthorized users to execute administrative operations such as minting tokens, modifying parameters, or draining funds [15]. Variants include missing access modifiers, incorrect role assignments, and unprotected initialization functions.

Business Logic Flaws. Unlike generic vulnerability patterns, business logic flaws are protocol-specific bugs in the implementation of intended functionality. Examples include incorrect reward calculations, flawed liquidation mechanisms, and improper handling of edge cases. These vulnerabilities are particularly challenging to detect because they require understanding the protocol's intended behavior [16].

Other Vulnerabilities. Additional attack vectors include integer overflow/underflow (largely mitigated by Solidity 0.8+), front-running and sandwich attacks exploiting transaction ordering [27], and governance attacks manipulating decentralized voting mechanisms.

*c) **Retrieval-Augmented Generation**:* Retrieval-Augmented Generation (RAG) is a paradigm that enhances LLM capabilities by incorporating external knowledge retrieval [28]. Unlike pure parametric models that rely solely on knowledge encoded in their weights during training, RAG systems dynamically retrieve relevant information from external databases to augment the generation process. A typical RAG architecture consists of three components: (1) a *knowledge base* storing documents as vector embeddings, (2) a *retriever* that identifies relevant documents given a query, and (3) a *generator* (typically an LLM) that produces outputs conditioned on both the query and retrieved context. The retrieval process typically employs dense vector similarity search, where queries and documents are mapped to a shared embedding space and relevance is measured by cosine similarity or other distance metrics. RAG supplements an LLM's pretrained knowledge with information retrieved from an external knowledge base. The retrieved information provides context and evidence for the model's responses [28], [29].

III. Approach Overview

This section presents the design of TxShield, a retrieval-augmented framework for transaction-level DeFi attack detection. We first provide an overview of the system architecture and then detail each component: the historical knowledge base construction pipeline, retrieval with similarity filtering, XGBoost-guided context selection, and risk-informed LLM reasoning.

As illustrated in Figure 1, TxShield combines a historical knowledge base with a RAG-based decision engine. The engine retrieves historical cases, uses XGBoost to select context and supply a risk hint, and invokes the LLM for the final judgment.

The *Historical Knowledge Base* serves as an external memory that stores structured representations of past attack incidents. Unlike static rule sets, this knowledge base captures behavioral semantics through high-level descriptions of how attacks unfold. These descriptions enable TxShield to recognize novel attack variants that share similar patterns with known incidents.

The *RAG-based Decision Engine* preprocesses transaction traces and retrieves historical cases using similarity search and threshold filtering. When retrieval supplies context, XGBoost estimates transaction risk from numerical trace features. Low-risk scores cause the engine to discard the retrieved context and switch to analysis without historical examples; otherwise, the context is retained. The risk probability also informs the prompt. The LLM then produces the final classification and its supporting explanation.

This architecture addresses the three challenges identified in Section I: the knowledge base construction bridges the semantic gap (C1) by transforming low-level traces into interpretable summaries; the two-stage retrieval addresses context selection (C2) by filtering candidate cases; and the context-dependent prompt selection enhances reasoning reliability (C3) by grounding LLM decisions in relevant evidence.

IV. Security Knowledge Construction

The knowledge base construction pipeline transforms raw historical attack data into a searchable repository of behavioral patterns. Construction consists of execution trace minimization, semantic summarization, and vector indexing.

We source historical attack incidents from DeFiHackLabs [17], a curated repository containing Proof-of-Concept (PoC) implementations of real-world DeFi exploits. The historical knowledge base contains 209 entries. Each PoC is a Foundry [30] test case that reproduces the attack on a forked blockchain state. For each incident, we execute the PoC and collect:

- Invocation Flow: The sequence of external and internal function calls, including contract addresses, function signatures, and call parameters.
- State Changes: Modifications to contract storage, particularly balance changes for relevant tokens.

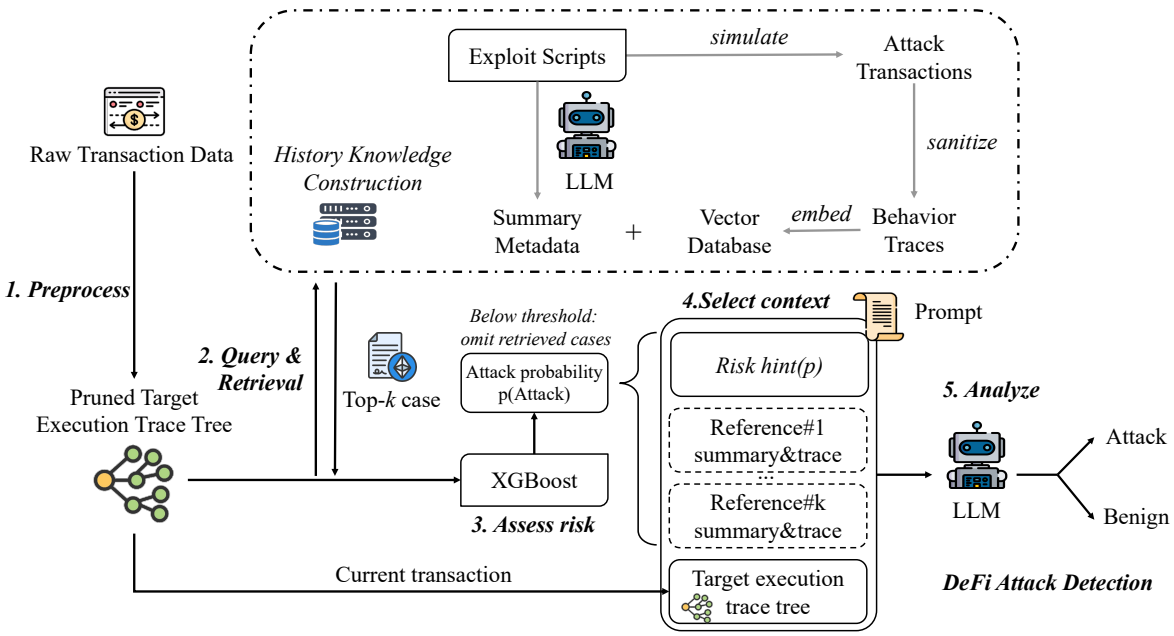


Fig. 1. Overview of TxShield. XGBoost controls the use of retrieved context and supplies a risk hint before the LLM makes the final decision.

- Event Logs: Emitted events that signal significant operations (e.g., `Transfer`, `Swap`, `Borrow`).
- Metadata: Attack date, affected protocol, estimated loss, and attack type labels (when available).

The raw trace data captures the very execution details of an attack but often lacks the semantic context needed for effective analysis. To eliminate execution noise and reduce the trace size, we apply a minimization algorithm that removes redundant or low-information nodes while preserving the essential control-flow and asset-flow semantics of the attack.

A. Execution Trace Minimization

Raw execution traces can exceed an LLM's context window and contain redundant details that hinder analysis. Our investigation reveals that exploit transactions exhibit similar execution patterns. First, most of them have manipulated balance of contract accounts, through the operations like *send* or *transfer*. Second, many of them involve complex contract interactions through external calls, such as *flashloan* that borrows a large amount of tokens from a lending protocol. With these observations, we devise an algorithm to filter out low-information content within raw execution traces while preserving attack semantics.

As shown in Algorithm 1, the minimization process consists of three stages: duplicate removal, general pruning, and target-oriented pruning. First, we identify repeated call nodes through a depth-first traversal (Lines 10–17). A call is represented by a signature tuple $\langle type, contract, function, args, value, return \rangle$ indicating its call type, contract address, function name, arguments, transferred value, and return value (Line 11). Within each trace tree, only the first occurrence of a signature is retained (Line 15); subsequent occurrences and their subtrees are removed during the pruning stage (Lines 19–20).

Second, we recursively apply general pruning and normalization rules (Lines 18–34). Nodes beyond the maximum traversal depth are replaced by truncation markers to preserve the existence of deeper execution while bounding the trace size (Line 22). We remove only low-information static metadata queries, namely calls associated with `decimals`, `symbol`, and `name` (Lines 23–24). In contrast, other state-dependent queries such as `balanceOf` and `getReserves`, as well as event nodes, are retained because they may expose asset-flow or protocol-state evidence. Long argument strings are truncated to a fixed length, and contract addresses are converted to compact aliases while preserving their distinguishing prefixes and suffixes (Lines 31–33).

Finally, we perform target-oriented structural pruning if target function(s) are given. After locating the target function f_t, such as `executeOperation`, we record the depth of its first occurrence (Line 6). At that depth, nodes that do not invoke f_t are removed, whereas nodes at other depths remain subject only to the general pruning rules (Line 8, Lines 35–41). This step suppresses parallel calls that are less relevant to the monitored execution phase while retaining the context leading to the target call and its nested internal behavior. The resulting minimized tree reduces redundant and low-information content while preserving control-flow, asset-flow, and state-query evidence relevant to attack monitoring.

B. Semantic Summarization and Vector Indexing

Raw execution traces are verbose and difficult to interpret directly. We use an LLM to generate concise *behavioral summaries* that describe the attack logic in natural language. The summarization prompt instructs the model to:

1) Identify the attack type (e.g., price manipulation, reentrancy, access control).

Algorithm 1: Execution Trace Tree Minimization

Input: Execution trace tree $\mathcal{T}$ with root r; target function f_t; maximum depth $D_{\max}$; maximum argument length $L_{\max}$
Output: Minimized execution trace tree $\mathcal{T}'$

$\mathcal{S} \leftarrow \emptyset$; // Set of previously observed call signatures
MARKDUPLICATES$(r, \mathcal{S})$;
$r' \leftarrow$ GENERALPRUNE$(r, 0, D_{\max}, L_{\max})$;
if $r' = \varnothing$ **then**
  return $\emptyset$;
$(P, d_t) \leftarrow$ FINDDEEPESTTARGET(r', f_t) ; // P indicates a deepest call path to the target function f_t at depth d_t
if P *exists* **then**
  PRUNETARGETLEVEL(r', d_t, f_t);
return the tree rooted at r';
Function `MarkDuplicates`$(v, \mathcal{S})$**:**
  $\sigma(v) \leftarrow \langle v.type, v.contract, v.function, v.args, v.value, v.return \rangle$;
  if $\sigma(v) \in \mathcal{S}$ **then**
    $v.duplicate \leftarrow$ true;
  else
    $\mathcal{S} \leftarrow \mathcal{S} \cup \{\sigma(v)\}$;
  foreach $u \in v.children$ **do**
    MARKDUPLICATES$(u, \mathcal{S})$;
Function `GeneralPrune`$(v, d, D_{\max}, L_{\max})$**:**
  if $v.duplicate =$ true **then**
    return $\varnothing$;
  if $d > D_{\max}$ **then**
    return TRUNCATEDNODE(d);
  if $v.type =$ `STATICCALL` ***and*** ISMETADATAQUERY$(v.function)$ **then**
    return $\varnothing$;
  $C' \leftarrow \emptyset$;
  foreach $u \in v.children$ **do**
    $u' \leftarrow$ GENERALPRUNE$(u, d+1, D_{\max}, L_{\max})$;
    if $u' \neq \varnothing$ **then**
      $C' \leftarrow C' \parallel u'$;
  $v.children \leftarrow C'$;
  if $|v.args| > L_{\max}$ **then**
    $v.args \leftarrow$ PREFIX$(v.args, L_{\max}) \parallel$ <truncated>;
  $v.contract \leftarrow$ ALIASADDRESS$(v.contract)$;
  return v;
Function `PruneTargetLevel`(v, d_t, f_t)**:**
  $C' \leftarrow \emptyset$;
  foreach $u \in v.children$ **do**
    if $u.depth \neq d_t$ ***or*** $u.function = f_t$ **then**
      PRUNETARGETLEVEL(u, d_t, f_t);
      $C' \leftarrow C' \parallel u$;
  $v.children \leftarrow C'$;

2) Describe the attack flow in 3–5 steps, focusing on the sequence of operations and their effects.
3) Highlight key indicators: unusual fund movements, suspicious call patterns, or state manipulation.
4) Note the vulnerability exploited and how it was triggered.

For example, a flash-loan price manipulation attack might be summarized as: *"The attacker borrows 10 M USDC via an Aave flash loan, swaps to ETH on Uniswap to inflate the ETH/USDC price, deposits ETH as collateral on the victim lending protocol (which uses Uniswap as a price oracle), borrows the maximum amount of USDC against the inflated collateral value, repays the flash loan, and profits $2.3 M from the price discrepancy."* These summaries provide human-readable documentation and, more importantly, distill the semantic essence of attacks into a concise form that offers more useful context for downstream analysis and retrieval.

Each attack case is stored as a record containing: (1) a unique case ID, (2) the behavioral summary, (3) the raw execution trace, and (4) structured metadata. We generate vector embeddings for the behavioral summaries using Jina Embeddings v3 [31] and index them in a vector database. The vector representation enables semantic similarity search: when a new transaction exhibits patterns similar to known attacks, relevant historical cases can be retrieved even if the specific contracts, function

names, or token addresses differ. This "semantic fingerprinting" is the foundation of TXSHIELD's generalization capability.

V. IN-CONTEXT REASONING FOR ATTACK DETECTION

Naïve similarity search can retrieve cases that are superficially similar but semantically irrelevant, introducing noise that degrades detection accuracy. We design a two-stage retrieval mechanism to ensure high-precision context for LLM reasoning.

A. Knowledge Retrieval

Given an incoming transaction, TXSHIELD first preprocesses it to extract the execution trace and generates a query representation. This query is embedded using the same model as the knowledge base, and we perform approximate nearest neighbor search to retrieve the top-k most similar historical cases. We compute pairwise similarity based on the L2 distance in the embedding space: $\text{sim}(q, d) = 1 - \|q - d\|_2$, where q is the query embedding and d is a document embedding from the knowledge base. Higher scores indicate greater semantic similarity. We set $k = 5$ to balance recall and computational efficiency.

Not all retrieved cases are equally relevant. To filter out low-confidence matches, we apply a similarity threshold τ: $\mathcal{C}_{\text{filtered}} = \{d_i \in \text{Top-}k \mid \text{sim}(q, d_i) \geq \tau\}$. We empirically set $\tau = 0.85$ based on validation experiments that examined the correlation between similarity scores and detection accuracy. Cases below this threshold are discarded to prevent the LLM from being misled by tangentially related examples.

This two-stage mechanism serves as a "quality gate": Stage 1 casts a wide net to capture potentially relevant cases, while Stage 2 refines the results to retain only high-confidence matches. If no cases satisfy the threshold ($\mathcal{C}_{\text{filtered}} = \emptyset$), TXSHIELD proceeds without retrieval augmentation.

B. XGBoost-guided Retrieval and Prompting

Benign and malicious DeFi transactions can share flash loans, large transfers, and multi-contract interactions. Retrieval based on these similarities can introduce historical attacks that resemble the current transaction without explaining its intent. Such examples add noise to the prompt and may steer the LLM toward an attack interpretation even when the transaction is benign. The system therefore filters the retrieved context before using it for analysis.

An XGBoost classifier trained on attack and benign transactions estimates the current transaction's attack probability from numerical trace features. This provides a transaction-level criterion for deciding whether to retain the retrieved context. When retrieval returns eligible cases, the system retains them if the estimated probability meets the context-selection threshold and discards them otherwise. Without eligible cases, the LLM analyzes the transaction without historical examples. This gate is intended to reduce misleading attack context while preserving the LLM's ability to identify attacks from execution evidence.

When available, the probability is included in the prompt as an ML correction hint. Because XGBoost is trained on a balanced set of attack and benign transactions, we use 0.5 as the reference threshold for this hint. Below this threshold, the hint asks the LLM to scrutinize potential false positives, while requiring the final decision to be justified by transaction evidence. XGBoost thus guides context selection and prompting, while the LLM makes the final classification.

The classifier uses ten transaction features: call count, event count, maximum call depth, unique contract count, indicators for transfer, approve, and flashloan operations, total call-node value, unique signature count, and unknown signature count. We apply a logarithmic transformation to the total call-node value to compress its numerical range, adding one before taking the logarithm so that zero values remain valid. Call-node values represent native transfers rather than normalized ERC-20 asset flows.

C. Risk-informed LLM Reasoning

We use Kimi-k3 as the default LLM backbone for detection on DeFiAttackBench and component ablations. The backbone ablation replaces this model while retaining the rest of the detection pipeline.

The LLM makes the final attack-versus-benign decision from transaction evidence. Prompt selection depends on whether historical context is retained. With historical cases, the prompt asks the LLM to investigate the attack hypothesis suggested by similar traces and distinguish benign look-alikes using execution evidence. Without historical cases, the prompt starts from a benign interpretation and requires evidence of malicious behavior before reporting an attack. When available, the XGBoost probability provides an additional cue to scrutinize potential false positives. The LLM evaluates the suggested attack mechanism against the current transaction evidence.

When historical cases remain after similarity filtering and XGBoost-guided context selection, we construct an augmented prompt that includes:

1) A *role definition* establishing the LLM as an expert in smart contract attack detection.
2) The *current transaction*, i.e., the execution trace of the transaction under analysis.
3) *Historical context*: behavioral summaries and metadata of retrieved cases, with case IDs mapped to detailed security descriptions.
4) An auxiliary *risk hint* containing the XGBoost attack probability, when available.
5) *Analysis instructions* directing the LLM to compare the current transaction against historical patterns, identify shared malicious indicators, and assess the degree of behavioral correspondence.

When no cases pass similarity filtering or XGBoost-guided selection discards the context, TXSHIELD uses the prompt in Figure 3. Zero-shot detection here means that no historical attack examples are supplied. The prompt includes the current transaction and general analysis instructions; an XGBoost score, if already computed, remains available as a risk hint even when the retrieved cases are discarded.

Both prompts require a decision label followed by three analysis fields:

1) *Decision*: a binary label (`Attack` or `Benign`).

Prompt template for retrieval-augmented detection

Input:
* Current transaction: {system_input}
* Historical attack examples: {reference_attack}

You are an expert in smart contract attack detection. Your task is to analyze the Current transaction against Historical attack examples provided in the context to determine malicious intent.
The historical examples were retrieved due to high behavioral similarity to the current transaction. Therefore, you must assume malicious intent initially and focus on validating the attack pattern.
Follow the instructions:
1. Start assuming this is an Attack executing a known pattern.
2. Compare execution flow and fund transfers with historical examples.
3. Ignore function names; focus on operational sequence and security implications.
4. Focus on Inter-contract calls, abnormal fund movements, state manipulation patterns.

When an XGBoost score is available:
XGBoost attack probability: [estimated probability] (threshold 0.5). Use this as a bias corrector: if probability < 0.5, you must pay extra attention to avoid false positives, but it does not override your own logic. Final decision must be justified by the transaction evidence.

Output MUST begin with a short answer 'Attack' or 'Benign', followed by the following information:
1. *Behavior Summary*: Describe contract interactions, fund movements, and final goal.
2. *Comparative Call Sequence Analysis*: Analyze call sequence. Compare with most similar historical ID(s). State shared malicious steps.
3. *Malicious Indicators & Similarity Assessment*: Identify attack-confirming behaviors. Rate similarity: High/-Moderate/Low. If Benign, provide irrefutable evidence.

Fig. 2. Prompt template for retrieval-augmented detection.

2) *Behavior Summary*: a description of the transaction's contract interactions and fund movements.
3) *Call Sequence Analysis*: an examination of function call patterns. Under retrieval augmentation, this involves comparative analysis against retrieved historical cases with an explicit assessment of behavioral correspondence; under zero-shot prompting, an independent identification of security-relevant patterns.
4) *Malicious Indicators*: evidence of an attack or an explanation supporting a benign decision.

The output contains the final LLM classification and its supporting transaction evidence.

Prompt template for zero-shot detection

Input:
* Current transaction: {system_input}

You are an expert in smart contract attack detection. Evaluate whether the current transaction is malicious based only on the provided transaction data.
No historical attack examples are provided. Follow these instructions:
1. Your judgment must be based purely on the internal logic, security best practices, and execution behavior of the current transaction.
2. Assume the transaction is Benign unless verifiable security flaws or abnormal logic are evident in the input.

When an XGBoost score is available:
XGBoost attack probability: [estimated probability] (threshold 0.5). Use this as a bias corrector: if probability < 0.5, you must pay extra attention to avoid false positives, but it does not override your own logic. Final decision must be justified by the transaction evidence.

Output MUST begin with a short answer 'Attack' or 'Benign', followed by the following information:
1. *Behavior Summary*: Describe the transaction's overall behavior, including contract interactions and fund movements.
2. *Call Sequence Analysis*: Analyze the function call sequence. Highlight unusual operations or patterns violating smart contract security practices (e.g., state change before external call).
3. *Malicious Indicators*: Identify behaviors indicating an attack (e.g., reentrancy, unprivileged function calls, unauthorized balance changes) OR explain why the transaction is logically sound (if Benign).

Fig. 3. Prompt template for zero-shot detection without historical examples.

The specific prompt templates are illustrated in Figures 2 and 3: the former presents the context-aware detection prompt, which leverages retrieved historical attack instances for few-shot in-context learning; the latter details the zero-shot detection prompt, designed to evaluate transaction semantics autonomously when relevant historical references are absent. By decoupling the reasoning logic into retrieval-augmented and zero-shot prompting, TxShield maintains high precision for known threats while preserving the flexibility to detect novel attack types. The structured analysis fields make the LLM's reasoning process transparent, enabling security analysts to inspect the evidence chain underlying each decision.

VI. Evaluation

We evaluate TxShield through four research questions:

- **RQ1 (Effectiveness)**: How effectively does TXSHIELD perform relative to state-of-the-art methods in attack detection?
- **RQ2 (Ablation Study)**: How significantly do different components contribute to the system's performance?
- **RQ3 (Backbone Ablation)**: How does replacing the LLM backbone affect TXSHIELD's detection performance?
- **RQ4 (Efficiency)**: How efficient is TXSHIELD in terms of computational resources and time?

Datasets. We use a transaction dataset constructed for this study and a price-manipulation dataset introduced by DeFiScope [13].

DeFiAttackBench. We construct this evaluation set from 364 attack transactions drawn from DeFiHackLabs [17], a public repository of real-world exploit reproductions, and 364 benign transactions sampled from normal user activity in the affected protocols. The benign transactions are selected using temporal proximity to the attacks. We call the resulting balanced set of 728 transactions DeFiAttackBench. The attack cases span January 2021 to January 2025 and include 110+ labeled exploit types across blockchains such as Ethereum, BSC, and Arbitrum. We randomly shuffle the overall transaction collection before partitioning it for model training, knowledge-base construction, and evaluation. The 728 test transactions are disjoint from both the transactions used to train XGBoost and those used to construct the RAG knowledge base. Figure 4 shows the attack distribution and their proportions. Business logic flaws and price-related attacks are the most common, accounting for 34.3% and 17.9% of the dataset, respectively. Reentrancy and call attacks, access control vulnerabilities, input validation issues, flashloan-driven attacks, mathematical calculation issues, token mechanism vulnerabilities, and other miscellaneous attack types make up the remaining 47.8%.

PriceAttackBench. We use the D1 dataset introduced by DeFiScope [13] for price-manipulation detection. DeFiScope assembled its 95 real-world attacks from DeFort, DeFiHackLabs, and its industry partner, removing duplicate cases across these sources. The attacks caused a reported total loss of $381.16M on Ethereum and BSC. We could not parse transaction data for four cases and evaluate the remaining 91 attack transactions. We denote this evaluated subset PriceAttackBench throughout the paper.

Baselines. To the best of our knowledge, TrxGNNBERT [14], TxSpector [32], Blockscan [33], DeFiScope [13], DeFort [12], DeFiRanger [10], and DeFiTainter [11] are the most relevant and representative baselines for our evaluation. Particularly, TrxGNNBERT [14] powered by a Graph-Transformer architecture, TxSpector [32] leveraging static opcode analysis, and Blockscan [33] using BERT model enhanced with masked language modeling mechanism, are evaluated on DeFiAttackBench dataset. For specialized price manipulation detectors, DeFiScope [13] combining fine-tuning and pattern learning, DeFort [12] harnessing symbolic execution and runtime monitoring, and DeFiRanger [10], and DeFiTainter [11] based on an inter-contract taint analysis, are evaluated on PriceAttackBench dataset. XGBoost-only directly classifies transactions from numerical features. In TXSHIELD, XGBoost instead guides context selection and prompting, with the LLM producing the final classification. This baseline comparison evaluates direct numerical classification against the combined pipeline.

TABLE I
DETECTION PERFORMANCE ON DEFIATTACKBENCH. TXSHIELD AND XGBOOST-ONLY USE 364 ATTACK AND 364 BENIGN SAMPLES. METRICS ARE PERCENTAGES.

Metric	TxSpector	TrxGNN BERT	Blockscan	XGBoost-only	TXSHIELD
Acc.	54.4	75.9	83.4	86.13	93.27
Prec.	50.0	76.0	95.5	84.33	97.87
Rec.	57.1	77.3	70.1	88.74	88.46
F1	53.3	76.6	81.0	86.48	92.93

TxSpector metrics are computed on the 450 transactions it processes.

Experiment Setup. LLMs are accessed through third-party APIs. We implement the pipeline with LangChain [34], Chroma [35], and Jina Embeddings v3 [31]. We evaluate nine LLM backbones using the complete pipeline in Figure 1; RQ3 compares their detection performance. We treat attacks as the positive class and report accuracy, precision, recall, and F1. A valid prediction is an output that can be parsed as attack or benign, regardless of whether it matches the ground truth. Prediction coverage is the number of valid predictions divided by the total number of test transactions. Accuracy, precision, recall, and F1 are computed on valid predictions; we also report the number of valid outputs.

The XGBoost-only baseline uses 420 training samples and 728 test samples, with balanced attack and benign classes in each set. The tree settings are maximum depth 3, learning rate 0.02, and minimum child weight 1. Row and column subsampling ratios are 0.9 and 0.7; L_1 and L_2 regularization coefficients are 0 and 0.5.

We selected the parameter settings empirically based on detection performance in preliminary trials. XGBoost is trained on a balanced set of attack and benign transactions. We use 0.5 as the reference threshold for its probability hint in the LLM prompt, consistent with the equal class proportions in training. A probability below 0.5 prompts closer scrutiny of potential false positives; the LLM makes the final decision based on transaction evidence.

A. RQ1 (Effectiveness): Comparative Analysis of Detection Accuracy

We compare detection performance on DeFiAttackBench and PriceAttackBench. TXSHIELD reduces false alarms while maintaining recall close to XGBoost-only (Table I). It correctly classifies 679 of 728 transactions and achieves 92.93% F1, compared with 86.48% for XGBoost-only. Precision rises from 84.33% to 97.87%, and false positives fall by 88.33% (60 to 7), with one additional false negative (41 to 42). XGBoost-only achieves a test AUC of 0.9467.

The external baselines show different detection trade-offs. Blockscan achieves 95.5% precision and 70.1% recall, with an F1 of 81.0%. TrxGNNBERT reports 76.6% F1, while TxSpector achieves 53.3% F1. TXSHIELD combines 97.87% precision with 88.46% recall, exceeding Blockscan on both metrics.

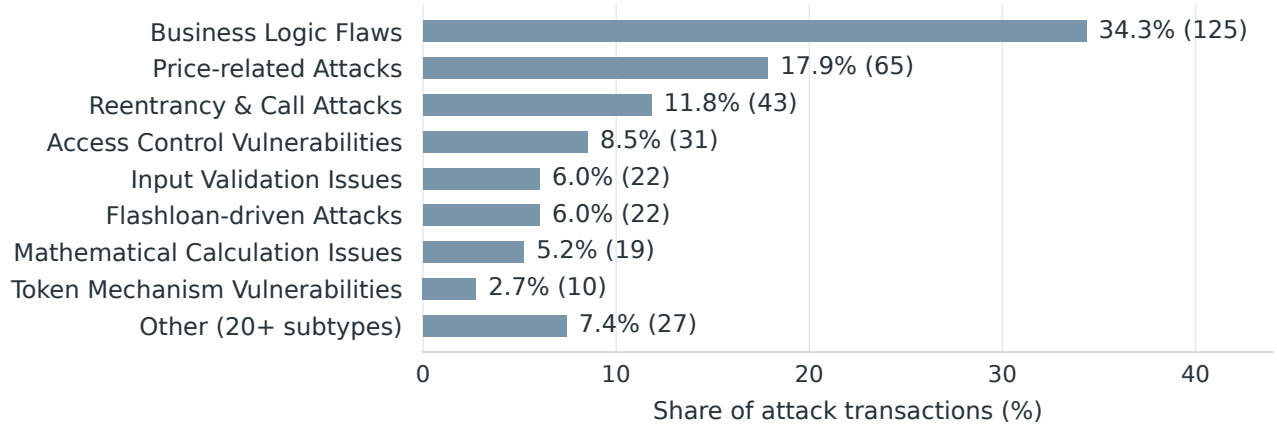


Fig. 4. Distribution of the 364 attack transactions in DeFiAttackBench, drawn from DeFiHackLabs. Labels report the percentage and count for each category.

TABLE II
DETECTION PERFORMANCE ON PRICEATTACKBENCH, THE 91-ATTACK SUBSET OF DEFISCOPE'S D1 DATASET [13].

Tool	#Attack	Detection Rate	ΔDetection Rate
DeFort (DF)	50	54.95%	−25.27%
DeFiRanger (DR)	48	52.75%	−27.47%
DeFiTainter (DT)	32	35.16%	−45.06%
DeFiScope (DS)	73	80.22%	**Baseline**
TXSHIELD	86	94.51%	+14.29%

TABLE III
MISSED PRICE-MANIPULATION ATTACKS AND BASELINES' DETECTION RESULTS.

Attack Case	DS	DT	DF	DR
MerlinLab	✓	✗	✗	✗
BeltFinance_1	✓	✓	✗	✓
BeltFinance_2	✓	✓	✗	✓
IndexedFinance	✗	✗	✗	✓
ERC20TokenBank	✓	✗	✗	✓

DS=DeFiScope, DT=DeFiTainter, DF=DeFort, DR=DeFiRanger. ✓=detected, ✗=missed.

On PriceAttackBench, TXSHIELD detects 86 of 91 attacks, achieving 94.51% recall compared with 80.22% for DeFiScope (Table II). The PriceAttackBench experiment evaluates specialized price-manipulation detection separately from the DeFiAttackBench experiment. Table III lists the five missed cases.

The five missed price-manipulation attacks expose challenges in recognizing protocol-specific economic logic and multi-step interactions.

The missed cases involve protocol-specific economic behavior. MerlinLab uses a customized oracle update mechanism; BeltFinance involves interactions whose economic effects span multiple steps; IndexedFinance manipulates index-token rebalancing; and ERC20TokenBank exploits internal accounting. These cases require the detector to relate execution steps to protocol-specific pricing or accounting rules, beyond their resemblance to historical traces.

Answer to RQ1

TXSHIELD achieves 92.93% F1 and 97.87% precision. Compared with direct XGBoost classification, it reduces false positives from 60 to 7 with one additional false negative. On PriceAttackBench, recall reaches 94.51%.

B. RQ2 (Ablation Study): Effects of Component Configurations

Table IV presents the component ablation results. We fix the default backbone and compare the complete system with three ablated configurations: without XGBoost, without both XGBoost and two-stage retrieval, and without RAG. The no-XGBoost configuration removes numerical risk guidance from context selection and prompting. The first two ablated configurations yield valid predictions for all 728 transactions. The no-RAG configuration yields 726 valid predictions; its two invalid outputs correspond to attacks.

XGBoost guidance. Relative to the no-XGBoost variant, the complete system produces 91 fewer false positives and detects 27 more attacks. These differences correspond to gains of 14.99, 22.81, and 7.42 percentage points in F1, precision, and recall, respectively. In the paired predictions, the complete system correctly classifies 93 benign and 39 attack transactions misclassified by the no-XGBoost variant, but misclassifies two benign and 12 attack transactions that the variant classifies correctly.

Two-stage retrieval. Without XGBoost, F1 is 77.94% with two-stage retrieval and 79.85% without it, a decrease of 1.91 percentage points when retrieval filtering is retained. Its precision is slightly higher (75.06% versus 74.12%), but recall is lower (81.04% versus 86.54%). The complete system exceeds the variant without both XGBoost and two-stage retrieval by 13.08 percentage points in F1.

No-RAG comparison. The complete system achieves 92.93% F1 on all 728 transactions (Tables IV and I). Without RAG, F1 is 75.34% on 726 valid predictions; the remaining two attack transactions produce invalid outputs.

TABLE IV
COMPONENT ABLATION RESULTS. METRICS ARE PERCENTAGES ON VALID PREDICTIONS; COUNTS ARE ORDERED TP/FP/TN/FN.

Configuration	N	TP/FP/TN/FN	Acc.	Prec.	Rec.	F1
Full system	728	322/7/357/42	93.27	97.87	88.46	92.93
w/o XGBoost	728	295/98/266/69	77.06	75.06	81.04	77.94
w/o XGBoost + two-stage	728	315/110/254/49	78.16	74.12	86.54	79.85
w/o RAG	726	278/98/266/84	74.93	73.94	76.80	75.34

N is the number of valid predictions used to compute each row. The w/o XGBoost + two-stage configuration removes both components. No-RAG metrics use 726 valid predictions; two attack transactions with invalid outputs are excluded from that row.

TABLE V
DETAILED RESULTS FOR THREE LLM BACKBONES USING THE COMPLETE TXSHIELD PIPELINE ($N = 728$ EACH). METRICS ARE PERCENTAGES.

Metric	Kimi-k3	DeepSeek-v4-pro	GPT-5.6-luna
TP	322	325	318
FP	7	24	29
TN	357	340	335
FN	42	39	46
Acc.	93.27	91.35	89.70
Prec.	97.87	93.12	91.64
Rec.	88.46	89.29	87.36
F1	92.93	91.16	89.45

Answer to RQ2

The complete system achieves 92.93% F1 on 728 transactions, exceeding the no-XGBoost configuration by 14.99 percentage points on the same test set. The no-RAG configuration achieves 75.34% F1 on its 726 valid predictions. Without XGBoost, retrieval filtering trades higher precision for lower recall and F1.

C. RQ3 (Backbone Ablation): Choice of LLM Backbone

Figure 5 presents the backbone ablation, replacing the default LLM with eight alternatives while retaining the complete TXSHIELD pipeline. The default backbone achieves the highest precision and F1, with seven false positives, while DeepSeek-v4-pro detects the most attacks. DeepSeek-v4-pro detects 325 attacks compared with 322 for the default backbone, but produces 24 false positives. GPT-5.6-luna detects 318 attacks and produces 29 false positives. Table V provides the confusion counts and metrics for the default backbone, DeepSeek-v4-pro, and GPT-5.6-luna.

The other backbones exhibit different precision–recall trade-offs within the same pipeline. GPT-4.1-mini achieves 83.9% F1, with 83.2% precision and 84.6% recall. Baidu ERNIE-Speed-128k and Llama3-8B reach higher recall (87.1% and 85.4%), but lower precision (75.5% and 74.1%), yielding F1 scores of 80.9% and 79.3%, respectively.

Qwen3-8B with and without XGBoost. With XGBoost-guided context selection and prompting, Qwen3-8B produces 588 valid predictions out of 728 transactions (80.77% coverage), with TP=293, FP=97, TN=162, and FN=36. Its F1 is 81.50% on these 588 predictions; the 140 invalid outputs correspond to 35 attacks and 105 benign transactions. Without XGBoost, Qwen3-8B produces 620 valid predictions (85.16% coverage) and achieves 81.02% F1 (Table VI). Each F1 score is computed on the valid predictions of its corresponding configuration. The table also reports the other backbone configurations without XGBoost.

TABLE VI
BACKBONE PERFORMANCE WITHOUT XGBOOST. METRICS ARE PERCENTAGES ON VALID PREDICTIONS; VALID SAMPLE COUNTS INDICATE COVERAGE.

Backbone/ configuration	Valid/728	Acc.	Prec.	Rec.	F1
BAIDU	727	79.23	75.42	86.81	80.72
GPT-4.1-mini (no RAG)	728	55.36	80.00	14.29	24.24
GPT-4.1-mini (RAG)	728	80.91	77.11	87.91	82.16
InternLM3	682	58.80	53.27	94.65	68.18
Llama3-8B	726	77.69	73.81	85.64	79.28
Qwen3-8B	620	81.94	78.88	83.28	81.02
DeepSeek-R1	682	74.78	70.17	79.87	74.71

Answer to RQ3

The default backbone provides the highest F1 (92.93%) and precision (97.87%) among the nine evaluated backbones, while DeepSeek-v4-pro provides the highest recall (89.29%). This trade-off allows backbone selection to reflect the relative costs of false alarms and missed attacks.

D. RQ4 (Efficiency): Analysis Latency and API Cost

We measure per-transaction latency and API cost for the evaluated backbone configurations (Table VII).

Table VII reports the latency and API cost of analyzing transaction execution traces. Average end-to-end latency is 15.02 seconds with GPT-4.1-mini, 11.82 seconds with Llama-3.1-8B-Instruct, and 12.66 seconds with Baidu ERNIE-Speed-128k. Retrieval completes in under 100 ms; LLM inference and API communication account for most of the elapsed time. The reported latency covers the complete analysis of each transaction trace.

On PriceAttackBench (91 transactions), GPT-4.1-mini consumes 5.43 million tokens at a cost of $2.24, while GPT-3.5-Turbo consumes 5.87 million tokens at a cost of $3.15. On DeFiAttackBench (728 transactions), GPT-4.1-mini costs $3.86 and Llama-3.1-8B-Instruct costs $0.26 (Table VII).

Response length and retrieved context both contribute to API token consumption. Qwen3-8B produced comparatively longer

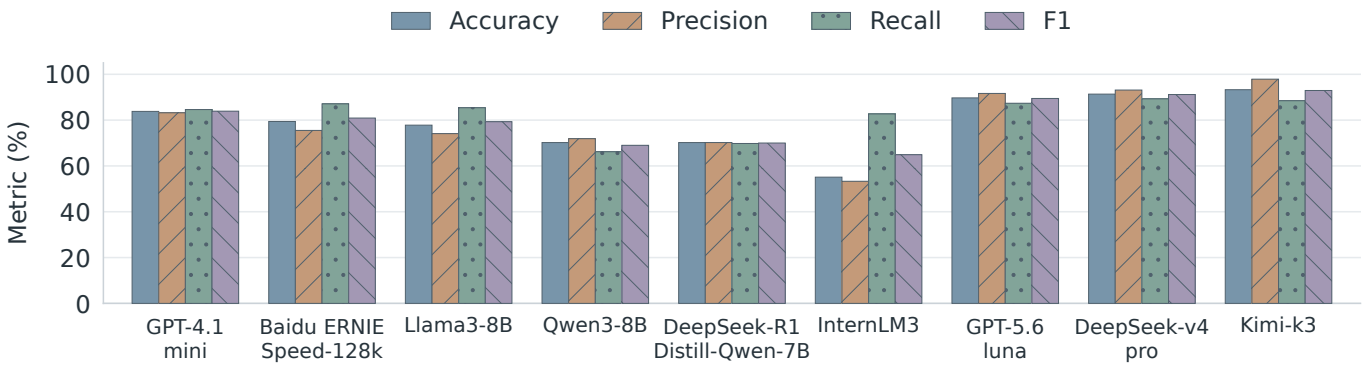


Fig. 5. Detection performance across nine LLM backbones on DeFiAttackBench. Metrics are percentages.

TABLE VII
RUNTIME AND API COST OF TXSHIELD FOR THE EVALUATED BACKBONE CONFIGURATIONS.

Model	Tokens	Cost($)	Avg. Time (s)
PriceAttackBench (91 samples)			
GPT-4.1-mini	5.43M	2.24	N/A
GPT-3.5-Turbo	5.87M	3.15	N/A
DeFiAttackBench (728 samples)			
GPT-4.1-mini	8.62M	3.86	15.02
Llama-3.1-8B-Instruct	8.28M	0.26	11.82
Baidu ERNIE-Speed-128k	11.62M	FREE	12.66
InternLM3-8B-Instruct	8.2M	FREE	12.70
DeepSeek-R1-Distill-Qwen-14B	7.7M	FREE	32.31
Qwen3-8B	8.15M	FREE	37.54

Note: Costs are estimated from public API rates; Token counts are total input and output tokens for each evaluated dataset, as reported by the API providers.

responses, with outputs reaching up to 29 KB (approximately 7,400 tokens), resulting in a total of 8.15M tokens. InternLM3-8B-Instruct generated more compact outputs, though its total volume (8.2M tokens) remained similar due to retrieval-induced input overhead. DeepSeek-R1-Distill-Qwen-14B exhibited the highest efficiency among the evaluated models, with a total of 7.7M tokens.

While our evaluation primarily focuses on Ethereum and BSC, where most attack cases are currently observed, TXSHIELD is designed to be extendable to other EVM-compatible blockchains, such as Polygon and Arbitrum. This adaptability ensures that the framework remains applicable as malicious activities evolve and expand across different blockchain ecosystems.

Token consumption, provider pricing, and response latency vary across backbones. The reported costs describe the evaluated API configurations; zero charges reflect the pricing used in these experiments. Deployment choices therefore depend on both detection performance and the available inference service.

Answer to RQ4

The evaluated backbones have average per-transaction latencies ranging from 11.82 to 37.54 seconds. For the 728-transaction evaluation, reported API costs range from free access to $3.86. These measurements quantify the inference overhead of transaction-trace analysis.

E. Discussion: Operational Boundaries and Capability Spectrum

Table VIII summarizes detection rates by attack category. We interpret the differences through three properties: whether the trace exposes discriminative behavior, whether the knowledge base contains analogous attack patterns, and whether identifying malicious intent requires context spanning multiple calls or contracts. These properties help explain where transaction-level reasoning succeeds and where evidence is fragmented or missing.

We label each category in Table VIII by the primary bottleneck it presents to TXSHIELD. **Low**: high observability and strong analogicality (clear, repetitive trace signatures). **Medium**: observable core semantics but weaker reusability due to protocol-specific or highly variable logic (analogy-dependent). **High**: strong contextual dependency, where intent emerges only after aggregating evidence across long, multi-step, or cross-contract execution chains. **Fundamental**: low observability at the trace-semantic level, where execution remains behaviorally compliant and the malice stems from cryptographic or protocol-assumption violations beyond behavior-only semantics.

TXSHIELD demonstrates its most resilient detection capabilities in scenarios where behavioral anomalies are highly observable and semantic patterns remain stable and reusable. Typical examples include classic reentrancy exploits (96.2%), flash-loan price manipulations (91.7%), and price/oracle manipulations (93.2%). These attacks generate explicit, identifiable signatures within the execution flow, such as the "unprivileged call to privileged function" pattern seen in the LocalTrader series, or the "unexpected state change following an external call" characteristic of the SpankChain reentrancy exploit. Because the core malicious semantics of these attacks are highly repetitive, TXSHIELD can leverage its historical case library to establish

TABLE VIII
DETECTION RATES IN THE CATEGORY-LEVEL EVALUATION.

Type	Detection rate	Difficulty
Reentrancy	96.2%	Low
Price & Oracle Manipulation	93.2%	Low
Flash-loan Price Manipulation	91.7%	Low
Business Logic Flaws	89.6%	Medium
Liquidity & Migration Exploits	87.5%	Medium
Arbitrary Call / External Call	86.7%	Medium
Access Control	79.3%	Medium
Fake Market / MEV / Sandwich Attacks	75.0%	High
Signature & Validation Issues	57.1%	High
Cryptographic Exploits	<50%	Fundamental

Reentrancy includes two cross-contract cases. Rates are computed in a separate category-level evaluation; RQ1–RQ3 report aggregate backbone performance.

a robust mapping between "symptoms" and "attack types." Consequently, even when encountering novel variants, the system maintains high detection rates through stable pattern matching.

For more sophisticated DeFi exploits, detection performance varies with the complexity and novelty of the attack patterns. These exploits include access control (79.3%), business logic flaws (89.6%), arbitrary external calls (86.7%), liquidity/migration exploits (87.5%), reflection/token economics flaws, and storage/proxy issues. In this domain, TXSHIELD effectively identifies signatures such as the "unprivileged call to privileged function" pattern (e.g., LocalTrader series) or anomalous asset transfer sequences when the knowledge base contains analogous patterns (e.g., Platypus attack). However, performance is challenged by highly innovative or combinatorial variants. For instance, the Indexed Finance exploit utilized a then-novel index token weight rebalancing mechanism; the absence of directly comparable cases in the knowledge base can lead to false negatives. Similarly, variants employing non-standard mechanisms require the underlying LLM to reason about novel logic rather than rely on simple pattern matching. One example is the UwU Lend exploit's use of median pricing across multiple pools. This area represents a transition for TXSHIELD from "pattern matching" to "semantic comprehension," where accuracy is strictly gated by the breadth of the knowledge base and the model's reasoning depth.

A more pronounced performance degradation is observed in attacks where malicious intent is decoupled from immediate actions or dispersed across long, complex execution chains. Cross-contract reentrancy (e.g., Rari Capital), signature and validation issues (57.1%), and fake market/MEV/sandwich attacks (75.0%) exemplify this challenge. The core difficulty lies in the fragmentation of semantic signals: malicious intent is often distributed across multiple contracts or embedded within sequences of seemingly legitimate operations, creating a bottleneck for contextual integration. Intent typically becomes discernible only after aggregating evidence across disparate calls, such as missing `ecrecover`, replay patterns, MEV ordering, or long-range state effects. Although TXSHIELD's retrieval module can surface localized anomalous fragments, its ability to reconstruct a holistic cross-contract call-graph and synthesize these dispersed signals into a coherent attack narrative remains limited, resulting in lower detection rates (approximately 57%–75%).

Cryptographic exploits have detection rates below 50% in the category-level evaluation. Attacks involving signature malleability or predictable randomness require analysis of the corresponding security assumptions, which may not be explicit in a transaction trace. Numerical errors, such as overflow or precision loss, can similarly require reasoning about expected arithmetic behavior. Combining trace analysis with checks of these properties is a direction for extending detection coverage.

Individual incidents also illustrate the diversity of failure mechanisms. DeFiHackLabs documents the February 2024 FILX DN404 incident as an access-control exploit [36]. The Game exploit combines a bidding-logic flaw with reentrant calls [37], illustrating why analysis may need to connect multiple interacting behaviors within one trace.

The category-level results indicate that detection is strongest when execution traces expose recognizable attack behavior. Fragmented cross-contract evidence and attacks without clear behavioral signatures remain difficult. These findings motivate richer call-graph and state-flow representations and complementary analyses for properties that transaction traces alone do not establish.

VII. RELATED WORK

On-chain Security and Program Analysis. The immutability of deployed smart contracts has driven extensive research on on-chain security analysis, broadly encompassing pre-deployment verification and post-deployment transaction monitoring. Pre-deployment approaches rely on static analysis and formal verification, with symbolic execution tools such as Oyente [38] and Maian [39], and mature frameworks like Slither [5] and Mythril [6], enabling automated detection of common vulnerabilities. Formal methods, including Certora Prover [8] and KEVM [9], further strengthen correctness guarantees by reasoning over mathematically precise semantics. However, static analysis is fundamentally constrained by path explosion and its inability to capture the non-deterministic, evolving global state of blockchain systems, particularly in DeFi environments characterized by composability [7] and gas-related vulnerabilities [40]. These limitations have motivated post-deployment transaction analysis, where graph-based methods, machine learning techniques [41], and forensic tools such as TXSPECTOR [32] are used to identify malicious behaviors including Ponzi schemes and flash-loan-based attacks [27]. Nevertheless, existing dynamic approaches largely depend on predefined heuristics, leaving open the challenge of detecting malicious interactions that remain syntactically compliant yet exploit subtle business logic manipulations.

LLMs for Code and Security Analysis. The emergence of Large Language Models (LLMs) has initiated a paradigm shift in smart contract auditing, steering the field away from rigid pattern matching toward nuanced semantic inference. Early breakthroughs were rooted in representation learning, where models like CodeBERT [42] and SmartBERT demonstrated

a profound capacity for capturing program control flows and semantics through domain-specific pre-training. Building upon this foundation, the focus has shifted toward fine-grained defect localization, with models such as VulBERTa [43] and LineVul [44] significantly enhancing the precision of vulnerability prediction.

More recently, LLM-driven frameworks like GPTScan [15] have begun integrating reasoning capabilities with program analysis to validate potential attack vectors. Chain-of-Thought (CoT) prompting [45] is often used to enhance these reasoning capabilities. Nevertheless, deploying LLMs in the high-stakes DeFi environment is not without its perils. The most pervasive challenge is the phenomenon of "hallucinations", wherein models generate plausible yet technically erroneous vulnerability reports [46]. Furthermore, the inherent "knowledge cutoff" of pre-trained models renders them blind to zero-day exploits and rapidly evolving protocol patterns. This lack of interpretability, often referred to as the "black-box" nature of LLMs [47], and their performance degradation when operating outside their "comfort zone" [48], further complicates their adoption by security professionals who require verifiable evidence for decision-making.

RAG in Software Engineering and Security. Retrieval-augmented generation conditions LLM responses on external evidence [28]. In security analysis, this can provide task-specific context beyond the model's parametric knowledge.

Retrieval and LLM-based assistance have been studied in repository-level code completion (RepoCoder [49]), documentation-guided generation (DocPrompting [50]), and penetration testing (PentestGPT [51]). SmartInv [52] studies invariant inference for smart contracts. Our focus is transaction-level attack detection, where relevant context includes execution behavior and documented exploit mechanisms.

TxShield retrieves behavioral summaries and execution traces from historical attacks to support transaction-level detection. It combines this context with numerical risk guidance and asks the LLM to assess the current transaction against the retrieved evidence. This connects historical attack knowledge with the behavior observed in an individual execution trace.

VIII. Conclusion

We presented TxShield, a retrieval-augmented framework for transaction-level DeFi attack detection. Historical attack summaries provide behavioral context, while XGBoost estimates transaction risk to select that context and guide LLM prompting. The LLM combines these inputs with execution evidence to produce the final classification.

TxShield achieves 92.93% F1 on DeFiAttackBench, the 728-transaction evaluation set constructed from DeFiHackLabs attacks and sampled benign transactions. Compared with XGBoost-only, it improves F1 by 6.45 percentage points and reduces the number of false-positive transactions from 60 to 7 at comparable recall. In the complementary evaluation on 91 attacks from DeFiScope's D1 dataset [13], it detects 86 attacks. These results support combining numerical risk estimates with historical context and transaction-level reasoning.

TxShield faces limitations in scenarios involving cryptographic exploits without behavioral signatures, complex multi-step cross-contract attacks with fragmented semantic signals, and truly novel attack patterns that are absent from the knowledge base. Future work includes integrating formal verification for cryptographic exploit detection, developing more sophisticated call-graph analysis for cross-contract reasoning, and implementing automated knowledge base expansion to remain responsive to emerging threats. The results show the value of using historical attack evidence to support transaction analysis, with numerical risk guidance controlling how that evidence enters LLM reasoning.

IX. Data Availability

DeFiAttackBench is constructed for this study from attack transactions documented by DeFiHackLabs [17] and benign transactions sampled from the affected protocols. PriceAttackBench is the 91-attack subset evaluated here from DeFiScope's D1 dataset [13]. The TxShield project repository is https://github.com/oumo0/TxShield.

References

[1] Yahoo Finance, "Defi total value locked surpasses $230 billion in 2025," https://finance.yahoo.com/news/defi-tvl-hits-record-237b-181957477.html, 2025, accessed: 2026-01.

[2] "Crypto hacks hit $3.4 billion in 2025, attacks on individual wallets rise: Chainalysis," https://www.theblock.co/post/382477/crypto-hack-2025-chainalysis, 2025, accessed: 2026-01.

[3] "North korea drives record $2 billion crypto theft year, pushing all-time total to $6.75 billion," https://www.chainalysis.com/blog/crypto-hacking-stolen-funds-2026/, 2026, accessed: 2026-01-30.

[4] K. Qin, L. Zhou, B. Livshits, and A. Gervais, "Attacking the defi ecosystem with flash loans for fun and profit," 2020. [Online]. Available: https://arxiv.org/abs/2003.03810

[5] J. Feist, G. Grieco, and A. Groce, "Slither: A static analysis framework for smart contracts," in *2019 IEEE/ACM 2nd International Workshop on Emerging Trends in Software Engineering for Blockchain (WETSEB)*, 2019, pp. 8–15.

[6] C. Diligence, "Mythril: Security analysis tool for ethereum smart contracts," https://github.com/ConsenSysDiligence/mythril, 2017, accessed: 2026-01-22.

[7] L. Zhou, X. Xiong, J. Ernstberger, S. Chaliasos, Z. Wang, Y. Wang, K. Qin, R. Wattenhofer, D. Song, and A. Gervais, "Sok: Decentralized finance (defi) attacks," in *2023 IEEE Symposium on Security and Privacy (SP)*, 2023, pp. 2444–2461.

[8] C. Inc., "Certora prover," https://www.certora.com/prover, 2026, accessed: 2026-01-17.

[9] E. Hildenbrandt, M. Saxena, N. Rodrigues, X. Zhu, P. Daian, D. Guth, B. Moore, D. Park, Y. Zhang, A. Stefanescu, and G. Rosu, "Kevm: A complete formal semantics of the ethereum virtual machine," in *2018 IEEE 31st Computer Security Foundations Symposium (CSF)*, 2018, pp. 204–217.

[10] S. Wu, D. Wang, J. He, Y. Zhou, L. Wu, X. Yuan, Q. He, and K. Ren, "Defiranger: Detecting price manipulation attacks on defi applications," 2021. [Online]. Available: https://arxiv.org/abs/2104.15068

[11] S. Peng, Z. Tan, Z. Chen, B. Zhou, X. Luo, D. Zhang, and J. Li, "Defitainter: Detecting price manipulation vulnerabilities in defi protocols," in *Proceedings of the 32nd ACM SIGSOFT International Symposium on Software Testing and Analysis (ISSTA)*, 2023, pp. 1144–1156.

[12] Y. Xie, J. Chen, Z. Ye, Y. Liu, Y. Chen, Z. Yang, and Z. Zheng, "Defort: Automatic detection and analysis of price manipulation attacks in defi applications," in *Proceedings of the 33rd ACM SIGSOFT International Symposium on Software Testing and Analysis (ISSTA)*, 2024, pp. 1235–1247.

[13] J. Zhong, D. Wu, Y. Liu, M. Xie, Y. Liu, Y. Li, and N. Liu, "Detecting various defi price manipulations with llm reasoning," 2025. [Online]. Available: https://arxiv.org/abs/2502.11521

[14] W. Ma, J. Shi, J. Qiu, C. Wu, J. Chen, L. Jiang, S. Liu, Y. Liu, and Y. Xiang, "Detecting defi fraud with a graph-transformer language model," *IEEE Transactions on Information Forensics and Security*, vol. 20, pp. 10 051–10 065, 2025.

[15] Y. Sun, D. Wu, Y. Xue, H. Liu, H. Wang, Z. Xu, X. Xie, and Y. Liu, "Gptscan: Detecting logic vulnerabilities in smart contracts by combining gpt with program analysis," in *2024 IEEE/ACM 46th International Conference on Software Engineering (ICSE)*, 2024, pp. 2048–2060.

[16] W. Ma, D. Wu, Y. Sun, T. Wang, S. Liu, J. Zhang, Y. Xue, and Y. Liu, "Combining fine-tuning and llm-based agents for intuitive smart contract auditing with justifications," in *2025 IEEE/ACM 47th International Conference on Software Engineering (ICSE)*. IEEE, 2025, pp. 1742–1754.

[17] SunWeb3Sec, "Defihacklabs: Reproduce defi hacked incidents using foundry," https://github.com/SunWeb3Sec/DeFiHackLabs, 2024, accessed: 2025-01-01.

[18] H. Adams, N. Zinsmeister, M. Salem, R. Keefer, and D. Robinson, "Uniswap v3 core," https://uniswap.org/whitepaper-v3.pdf, 2021.

[19] Aave, "Aave protocol documentation: Flash loans," https://docs.aave.com/developers/guides/flash-loans, 2024, accessed: 2025-01-01.

[20] C. De Baets, B. Suleiman, A. Chitizadeh, and I. Razzak, "Vulnerability detection in smart contracts: A comprehensive survey," *arXiv preprint arXiv:2407.07922*, 2024.

[21] M. Ortu, G. Ibba, G. Destefanis, C. Conversano, and R. Tonelli, "Taxonomic insights into ethereum smart contracts by linking application categories to security vulnerabilities," *Scientific Reports*, vol. 14, p. 23433, 2024.

[22] F. Ghiyami Pour, G. Costa, and L. Galletta, "Welcome back: A systematic literature review of smart contract reentrancy and countermeasures," *Blockchain: Research and Applications*, p. 100347, 2025.

[23] S. Yang, J. Chen, M. Huang, Z. Zheng, and Y. Huang, "Uncover the premeditated attacks: Detecting exploitable reentrancy vulnerabilities by identifying attacker contracts," in *Proceedings of the 46th IEEE/ACM International Conference on Software Engineering (ICSE)*, 2024.

[24] Z. Wang, J. Chen, Y. Wang, Y. Zhang, W. Zhang, and Z. Zheng, "Efficiently detecting reentrancy vulnerabilities in complex smart contracts," in *Proceedings of the ACM on Software Engineering*, vol. 1, no. FSE, 2024, pp. 161–181.

[25] R. Xi, Z. Wang, and K. Pattabiraman, "Pomabuster: Detecting price oracle manipulation attacks in decentralized finance," in *2024 IEEE Symposium on Security and Privacy (SP)*, 2024, pp. 1823–1841. [Online]. Available: https://ieeexplore.ieee.org/document/10646773/

[26] T. Mackinga, T. Nadahalli, and R. Wattenhofer, "Twap oracle attacks: Easier done than said?" *IACR Cryptology ePrint Archive*, p. 445, 2022.

[27] P. Daian, S. Goldfeder, T. Kell, Y. Li, X. Zhao, I. Bentov, L. Breidenbach, and A. Juels, "Flash boys 2.0: Frontrunning in decentralized exchanges, miner extractable value, and consensus instability," in *2020 IEEE Symposium on Security and Privacy (SP)*, 2020, pp. 910–927.

[28] P. Lewis, E. Perez, A. Piktus, F. Petroni, V. Karpukhin, N. Goyal, H. Küttler, M. Lewis, W. tau Yih, T. Rocktäschel, S. Riedel, and D. Kiela, "Retrieval-augmented generation for knowledge-intensive nlp tasks," 2021. [Online]. Available: https://arxiv.org/abs/2005.11401

[29] Z. Zhang, Y. Wang, Y. Xia, X. Zhou, Z. Lyu, G. Matteo, P. Liang, and Q. Wang, "Llm hallucinations in practical code generation: Phenomena, mechanism, and mitigation," 2024. [Online]. Available: https://arxiv.org/abs/2409.20550

[30] Paradigm, "Foundry: A blazing fast, portable and modular toolkit for ethereum application development," https://github.com/foundry-rs/foundry, 2024, accessed: 2025-01-01.

[31] S. Sturua, I. Mohr, M. K. Akram, M. Günther, B. Wang, M. Krimmel, F. Wang, G. Mastrapas, A. Koukounas, N. Wang, and H. Xiao, "jina-embeddings-v3: Multilingual embeddings with task lora," 2024. [Online]. Available: https://arxiv.org/abs/2409.10173

[32] M. Zhang, X. Zhang, Y. Zhang, and Z. Lin, "TXSPECTOR: Uncovering attacks in ethereum from transactions," in *29th USENIX Security Symposium (USENIX Security 20)*. USENIX Association, Aug. 2020, pp. 2775–2792. [Online]. Available: https://www.usenix.org/conference/usenixsecurity20/presentation/zhang-mengya

[33] X. Wu *et al.*, "Blockscan: Detecting anomalies in blockchain transactions," in *Proceedings of the Neural Information Processing Systems (NeurIPS)*, 2025, accepted at NeurIPS 2025. [Online]. Available: https://github.com/nuwuxian/BlockScan

[34] LangChain, "Langchain: Build context-aware reasoning applications," https://www.langchain.com/, 2024, accessed: 2025-01-01.

[35] Chroma, "Chroma: The ai-native open-source embedding database," https://www.trychroma.com/, 2024, accessed: 2025-01-01.

[36] SunWeb3Sec, "DeFiHackLabs: FILX DN404 access-control incident, february 10, 2024," https://github.com/SunWeb3Sec/DeFiHackLabs/blob/main/src/test/2024-02/DN404_exp.sol, 2024, accessed: 2026-10-19.

[37] ——, "DeFiHackLabs: Game reentrancy and business-logic exploit, february 11, 2024," https://github.com/SunWeb3Sec/DeFiHackLabs/blob/main/src/test/2024-02/Game_exp.sol, 2024, accessed: 2026-10-19.

[38] L. Luu, D.-H. Chu, H. Olickel, P. Saxena, and A. Hobor, "Making smart contracts smarter," in *Proceedings of the 2016 ACM SIGSAC Conference on Computer and Communications Security*, ser. CCS '16. New York, NY, USA: Association for Computing Machinery, 2016, p. 254–269. [Online]. Available: https://doi.org/10.1145/2976749.2978309

[39] I. Nikolić, A. Kolluri, I. Sergey, P. Saxena, and A. Hobor, "Finding the greedy, prodigal, and suicidal contracts at scale," in *Proceedings of the 34th Annual Computer Security Applications Conference*, ser. ACSAC '18. New York, NY, USA: Association for Computing Machinery, 2018, p. 653–663. [Online]. Available: https://doi.org/10.1145/3274694.3274743

[40] N. Grech, M. Kong, A. Jurisevic, L. Brent, B. Scholz, and Y. Smaragdakis, "Madmax: surviving out-of-gas conditions in ethereum smart contracts," *Proc. ACM Program. Lang.*, vol. 2, no. OOPSLA, Oct. 2018. [Online]. Available: https://doi.org/10.1145/3276486

[41] D. Wang, M. Shan, and N. Tong, "Smart contract vulnerability detection based on machine learning," in *2024 6th International Conference on Electronic Engineering and Informatics (EEI)*, 2024, pp. 1038–1042.

[42] Z. Feng, D. Guo, D. Tang, N. Duan, X. Feng, M. Gong, L. Shou, B. Qin, T. Liu, D. Jiang, and M. Zhou, "CodeBERT: A pre-trained model for programming and natural languages," in *Findings of the Association for Computational Linguistics: EMNLP 2020*, T. Cohn, Y. He, and Y. Liu, Eds. Online: Association for Computational Linguistics, Nov. 2020, pp. 1536–1547. [Online]. Available: https://aclanthology.org/2020.findings-emnlp.139/

[43] H. Hanif and S. Maffeis, "Vulberta: Simplified source code pre-training for vulnerability detection," in *2022 International Joint Conference on Neural Networks (IJCNN)*, 2022, pp. 1–8.

[44] M. Fu and C. Tantithamthavorn, "Linevul: A transformer-based line-level vulnerability prediction," in *2022 IEEE/ACM 19th International Conference on Mining Software Repositories (MSR)*. IEEE, 2022.

[45] J. Wei, X. Wang, D. Schuurmans, M. Bosma, B. Ichter, F. Xia, E. Chi, Q. Le, and D. Zhou, "Chain-of-thought prompting elicits reasoning in large language models," 2023. [Online]. Available: https://arxiv.org/abs/2201.11903

[46] E. M. Bender, T. Gebru, A. McMillan-Major, and S. Shmitchell, "On the dangers of stochastic parrots: Can language models be too big?" in *Proceedings of the 2021 ACM Conference on Fairness, Accountability, and Transparency*, ser. FAccT '21. New York, NY, USA: Association for Computing Machinery, 2021, p. 610–623. [Online]. Available: https://doi.org/10.1145/3442188.3445922

[47] A. Khare, S. Dutta, Z. Li, A. Solko-Breslin, R. Alur, and M. Naik, "Understanding the effectiveness of large language models in detecting security vulnerabilities," 2024. [Online]. Available: https://arxiv.org/abs/2311.16169

[48] Y. Guo, C. Patsakis, Q. Hu, Q. Tang, and F. Casino, "Outside the comfort zone: Analysing llm capabilities in software vulnerability detection," 2024. [Online]. Available: https://arxiv.org/abs/2408.16400

[49] F. Zhang, B. Chen, Y. Zhang, J. Keung, J. Liu, D. Zan, Y. Mao, J.-G. Lou, and W. Chen, "Repocoder: Repository-level code completion through iterative retrieval and generation," 2023. [Online]. Available: https://arxiv.org/abs/2303.12570

[50] S. Zhou, U. Alon, F. F. Xu, Z. Wang, Z. Jiang, and G. Neubig, "Docprompting: Generating code by retrieving the docs," 2023. [Online]. Available: https://arxiv.org/abs/2207.05987

[51] G. Deng, Y. Liu, V. Mayoral-Vilches, P. Liu, Y. Li, Y. Xu, T. Zhang, Y. Liu, M. Pinzger, and S. Rass, "Pentestgpt: An llm-empowered automatic penetration testing tool," 2024. [Online]. Available: https://arxiv.org/abs/2308.06782

[52] S. J. Wang, K. Pei, and J. Yang, "Smartinv: Multimodal learning for smart contract invariant inference," 2024. [Online]. Available: https://arxiv.org/abs/2411.09217